

Experiment No. 3
Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:7/8/23
Date of Submission: 14/8/23

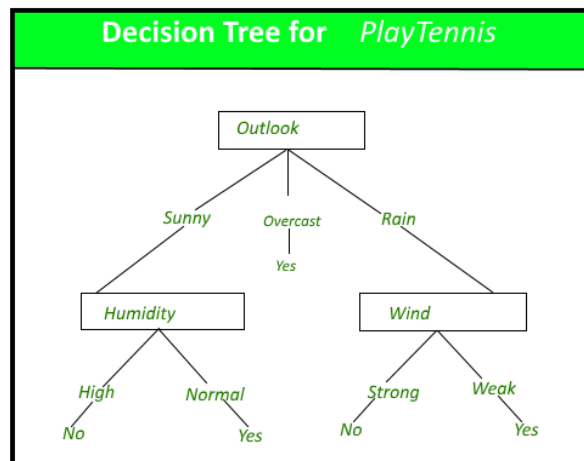
**Aim:** Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model.

**Objective:** To perform various feature engineering tasks, apply Decision Tree Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score. Improve the performance by performing different data engineering and feature engineering tasks.

**Theory:**



Decision Tree is the most powerful and popular tool for classification and prediction. A Decision tree is a flowchart-like tree structure, where each internal node denotes a test on an attribute, each branch represents an outcome of the test, and each leaf node (terminal node) holds a class label.



### **Dataset:**

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.



## Vidyavardhini's College of Engineering & Technology

### Department of Computer Engineering

---

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad Tobago, Peru, Hong, Holand-Netherlands.

**Code:**

```
# Import libraries
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

# To ignore warning messages
import warnings
warnings.filterwarnings('ignore')

# Adult dataset path
adult_dataset_path = "/content/adult.csv"

# Function for loading adult dataset
def load_adult_data(adult_path=adult_dataset_path):
    csv_path = os.path.join(adult_path)
    return pd.read_csv(csv_path)

# Calling load adult function and assigning to a new variable df
df = load_adult_data()
# load top 3 rows values from adult dataset
df.head(3)
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gain	capital.loss	hours.per.week	native.country	income
0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family	White	Female	0	0	40	USA	2156
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female	0	0	40	USA	2156
2	66	?	198661	Some-	10	Widowed	?	Unmarried	Black	Female	0	0	40	USA	2156

```
print ("Rows      : ",df.shape[0])
print ("Columns   : ",df.shape[1])
print ("\nFeatures : \n",df.columns.tolist())
print ("\nMissing values : ", df.isnull().sum().values.sum())
print ("\nUnique values : \n",df.nunique())
```

```
Rows      : 32561
Columns   : 15
```

```
Features :
['age', 'workclass', 'fnlwgt', 'education', 'education.num', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'capital.gain', 'capital.loss', 'hours.per.week', 'native.country', 'income']
```

```
Missing values : 0
```

```
Unique values :
age                73
workclass           9
fnlwgt             21648
education           16
education.num       16
marital.status       7
occupation          15
relationship         6
race                 5
sex                  2
capital.gain         119
capital.loss         92
hours.per.week       94
native.country       42
income               2
dtype: int64
```

```
# Let's understand the type of values present in each column of our adult dataframe 'df'.
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt                32561 non-null  int64
3   education             32561 non-null  object
4   education.num         32561 non-null  int64
5   marital.status        32561 non-null  object
6   occupation            32561 non-null  object
7   relationship          32561 non-null  object
```

```

8  race          32561 non-null object
9  sex           32561 non-null object
10 capital.gain  32561 non-null int64
11 capital.loss  32561 non-null int64
12 hours.per.week 32561 non-null int64
13 native.country 32561 non-null object
14 income        32561 non-null object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB

```

```

# Numerical feature of summary/description
df.describe()

```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week
<b>count</b>	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
<b>mean</b>	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40.437456
<b>std</b>	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	12.347429
<b>min</b>	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
<b>25%</b>	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40.000000
<b>50%</b>	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
<b>75%</b>	48.000000	2.370510e+05	12.000000	0.000000	0.000000	45.000000
<b>max</b>	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

```

# pull top 5 row values to understand the data and how it's look like
df.head()

```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	Fe
<b>0</b>	90	?	77053	HS-grad	9	Widowed	?	Not-in-family	White	Fe
<b>1</b>	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Fe
<b>2</b>	66	?	186061	Some-college	10	Widowed	?	Unmarried	Black	Fe
<b>3</b>	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Fe

```

# checking "?" total values present in particular 'workclass' feature
df_check_missing_workclass = (df['workclass']=='?').sum()
df_check_missing_workclass

```

1836

```

# checking "?" total values present in particular 'occupation' feature
df_check_missing_occupation = (df['occupation']=='?').sum()
df_check_missing_occupation

```

1843

```

# checking "?" values, how many are there in the whole dataset
df_missing = (df=='?').sum()
df_missing

```

```

age          0
workclass    1836
fnlwgt       0
education    0
education.num 0
marital.status 0
occupation   1843
relationship  0
race         0
sex          0
capital.gain  0
capital.loss  0
hours.per.week 0
native.country 583
income       0
dtype: int64

```

```

percent_missing = (df=='?').sum() * 100/len(df)
percent_missing

```

```

age          0.000000
workclass    5.638647

```

```

fnlwgt      0.000000
education   0.000000
education.num 0.000000
marital.status 0.000000
occupation   5.660146
relationship 0.000000
race         0.000000
sex          0.000000
capital.gain 0.000000
capital.loss 0.000000
hours.per.week 0.000000
native.country 1.790486
income      0.000000
dtype: float64

```

```

# Let's find total number of rows which doesn't contain any missing value as '?'
df.apply(lambda x: x != '?', axis=1).sum()

```

```

age          32561
workclass    30725
fnlwgt       32561
education    32561
education.num 32561
marital.status 32561
occupation   30718
relationship  32561
race         32561
sex          32561
capital.gain  32561
capital.loss  32561
hours.per.week 32561
native.country 31978
income       32561
dtype: int64

```

```

# dropping the rows having missing values in workclass
df = df[df['workclass'] != '?']
df.head()

```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Female
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female
5	34	Private	240000	HS-grad	9	Divorced	Other-service	Unmarried	White	Female

```

# select all categorical variables
df_categorical = df.select_dtypes(include=['object'])

# checking whether any other column contains '?' value
df_categorical.apply(lambda x: x == '?', axis=1).sum()

```

```

workclass      0
education      0
marital.status 0
occupation      7
relationship    0
race           0
sex            0
native.country 556
income         0
dtype: int64

```

```

# dropping the "?"s from occupation and native.country
df = df[df['occupation'] != '?']
df = df[df['native.country'] != '?']

```

```

# check the dataset whether cleaned or not?
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 1 to 32560
Data columns (total 15 columns):
#   Column          Non-Null Count  Dtype
---  -
0   age             30162 non-null  int64
1   workclass       30162 non-null  object
2   fnlwgt          30162 non-null  int64
3   education       30162 non-null  object

```

```
4 education.num 30162 non-null int64
5 marital.status 30162 non-null object
6 occupation 30162 non-null object
7 relationship 30162 non-null object
8 race 30162 non-null object
9 sex 30162 non-null object
10 capital.gain 30162 non-null int64
11 capital.loss 30162 non-null int64
12 hours.per.week 30162 non-null int64
13 native.country 30162 non-null object
14 income 30162 non-null object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
from sklearn import preprocessing
```

```
# encode categorical variables using label Encoder
```

```
# select all categorical variables
df_categorical = df.select_dtypes(include=['object'])
df_categorical.head()
```

	workclass	education	marital.status	occupation	relationship	race	sex	native.country	income
1	Private	HS-grad	Widowed	Exec-managerial	Not-in-family	White	Female	United-States	<=50k
3	Private	7th-8th	Divorced	Machine-op-inspct	Unmarried	White	Female	United-States	<=50k
4	Private	Some-college	Separated	Prof-specialty	Own-child	White	Female	United-States	<=50k
5	Private	HS-grad	Single	Other-service	Not-in-family	White	Female	United-States	<=50k

```
# apply label encoder to df_categorical
le = preprocessing.LabelEncoder()
df_categorical = df_categorical.apply(le.fit_transform)
df_categorical.head()
```

	workclass	education	marital.status	occupation	relationship	race	sex	native.country	income
1	2	11	6	3	1	4	0	38	0
3	2	5	0	6	4	4	0	38	0
4	2	15	5	9	3	4	0	38	0
5	2	11	0	7	4	4	0	38	0
6	2	0	5	0	4	4	1	38	0

```
# Next, Concatenate df_categorical dataframe with original df (dataframe)
```

```
# first, Drop earlier duplicate columns which had categorical values
df = df.drop(df_categorical.columns,axis=1)
df = pd.concat([df,df_categorical],axis=1)
df.head()
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week	workclass	education	marital
1	82	132870	9	0	4356	18	2	11	
3	54	140359	4	0	3900	40	2	5	
4	41	264663	10	0	3900	40	2	15	
5	34	216864	9	0	3770	45	2	11	
6	38	150601	6	0	3770	40	2	0	

```
# look at column type
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 1 to 32560
Data columns (total 15 columns):
#   Column              Non-Null Count  Dtype
---  -
0   age                  30162 non-null  int64
1   fnlwgt               30162 non-null  int64
2   education.num        30162 non-null  int64
3   capital.gain         30162 non-null  int64
4   capital.loss         30162 non-null  int64
5   hours.per.week       30162 non-null  int64
6   workclass            30162 non-null  int64
```

```

7   education      30162 non-null  int64
8   marital.status 30162 non-null  int64
9   occupation     30162 non-null  int64
10  relationship   30162 non-null  int64
11  race           30162 non-null  int64
12  sex            30162 non-null  int64
13  native.country 30162 non-null  int64
14  income         30162 non-null  int64
dtypes: int64(15)
memory usage: 3.7 MB

```

```

# convert target variable income to categorical
df['income'] = df['income'].astype('category')

```

```

# check df info again whether everything is in right format or not
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 1 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   30162 non-null  int64
1   fnlwgt                30162 non-null  int64
2   education.num         30162 non-null  int64
3   capital.gain          30162 non-null  int64
4   capital.loss          30162 non-null  int64
5   hours.per.week        30162 non-null  int64
6   workclass             30162 non-null  int64
7   education             30162 non-null  int64
8   marital.status        30162 non-null  int64
9   occupation            30162 non-null  int64
10  relationship          30162 non-null  int64
11  race                  30162 non-null  int64
12  sex                   30162 non-null  int64
13  native.country        30162 non-null  int64
14  income                 30162 non-null  category
dtypes: category(1), int64(14)
memory usage: 3.5 MB

```

```

# Importing train_test_split
from sklearn.model_selection import train_test_split

```

```

# Putting independent variables/features to X
X = df.drop('income',axis=1)

```

```

# Putting response/dependent variable/feature to y
y = df['income']

```

```
X.head(3)
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week	workclass	education	marita
1	82	132870	9	0	4356	18	2	11	
3	54	140359	4	0	3900	40	2	5	
4	41	264663	10	0	3900	40	2	15	

```
y.head(3)
```

```

1    0
3    0
4    0
Name: income, dtype: category
Categories (2, int64): [0, 1]

```

```

# Splitting the data into train and test
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.30,random_state=99)

```

```
X_train.head()
```



```

    age  fnlwgt  education.num  capital.gain  capital.loss  hours.per.week  workclass  education  ma
24351   42  289636             9             0             0             46           2           11
15626   37   52465             9             0             0             40           1           11

# Importing decision tree classifier from sklearn library
from sklearn.tree import DecisionTreeClassifier

# Fitting the decision tree with default hyperparameters, apart from
# max_depth which is 5 so that we can plot and read the tree.
dt_default = DecisionTreeClassifier(max_depth=5)
dt_default.fit(X_train,y_train)

DecisionTreeClassifier
DecisionTreeClassifier(max_depth=5)

# Let's check the evaluation metrics of our default model

# Importing classification report and confusion matrix from sklearn metrics
from sklearn.metrics import classification_report,confusion_matrix,accuracy_score

# making predictions
y_pred_default = dt_default.predict(X_test)

# Printing classifier report after prediction
print(classification_report(y_test,y_pred_default))

      precision    recall  f1-score   support

     0       0.86      0.95      0.91      6867
     1       0.78      0.52      0.63      2182

 accuracy      0.82
 macro avg      0.82
weighted avg      0.84

# Printing confusion matrix and accuracy
print(confusion_matrix(y_test,y_pred_default))
print(accuracy_score(y_test,y_pred_default))

[[6553  314]
 [1039 1143]]
0.8504807161012267

pip install pydotplus

Requirement already satisfied: pydotplus in /usr/local/lib/python3.10/dist-packages (2.0.2)
Requirement already satisfied: pyparsing>=2.0.1 in /usr/local/lib/python3.10/dist-packages (from pydotplus) (3.1.1)

# Importing required packages for visualization
from IPython.display import Image
from sklearn.externals.six import StringIO
from sklearn.tree import export_graphviz
import pydotplus,graphviz

# Putting features
features = list(df.columns[1:])
features


['fnlwgt',
 'education.num',
 'capital.gain',
 'capital.loss',
 'hours.per.week',
 'workclass',
 'education',
 'marital.status',
 'occupation',
 'relationship',
 'race',
 'sex',
 'native.country',
 'income']

import six
import sys
sys.modules['sklearn.externals.six'] = six

```

```
# plotting tree with max_depth=3
dot_data = StringIO()
export_graphviz(dt_default, out_file=dot_data,
                feature_names=features, filled=True, rounded=True)
```

```
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```




```
# GridSearchCV to find optimal max_depth
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
```

```
# specify number of folds for k-fold CV
n_folds = 5
```

```
# parameters to build the model on
parameters = {'max_depth': range(1, 40)}
```

```
# instantiate the model
dtree = DecisionTreeClassifier(criterion = "gini",
                              random_state = 100)
```

```
# fit tree on training data
tree = GridSearchCV(dtree, parameters,
                    cv=n_folds,
                    scoring="accuracy")
tree.fit(X_train, y_train)
```



```
# scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_max_depth	params	split0_t
0	0.014924	0.004959	0.004501	0.002179	1	{'max_depth': 1}	
1	0.019653	0.000594	0.003418	0.000383	2	{'max_depth': 2}	
2	0.025445	0.000331	0.003681	0.000198	3	{'max_depth': 3}	
3	0.031828	0.000514	0.003868	0.000580	4	{'max_depth': 4}	
4	0.043398	0.007946	0.003806	0.001610	5	{'max_depth': 5}	

```
"""
# plotting accuracies with max_depth
plt.figure()
plt.plot(scores["param_max_depth"],
         scores["mean_train_score"],
         label="training accuracy")
plt.plot(scores["param_max_depth"],
```

```

        scores["mean_test_score"],
        label="test accuracy")
plt.xlabel("max_depth")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
"""

'\n# plotting accuracies with max_depth\nplt.figure()\nplt.plot(scores["param_max_depth"], \n
scores["mean_train_score"], \n          label="training accuracy")\nplt.plot(scores["param_max_depth"],
\n          scores["mean_test_score"], \n          label="test accuracy")\nplt.xlabel("max_depth")\nplt.

# GridSearchCV to find optimal max_depth
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV


# specify number of folds for k-fold CV
n_folds = 5

# parameters to build the model on
parameters = {'min_samples_leaf': range(5, 200, 20)}

# instantiate the model
dtree = DecisionTreeClassifier(criterion = "gini",
                              random_state = 100)

# fit tree on training data
tree = GridSearchCV(dtree, parameters,
                    cv=n_folds,
                    scoring="accuracy")
tree.fit(X_train, y_train)

```



```

# scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()

```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_min_samples_leaf	par
0	0.071311	0.002583	0.002760	0.000105	5	{'min_samples_le
1	0.058933	0.001238	0.002700	0.000055	25	{'min_samples_le
2	0.054030	0.001434	0.002561	0.000125	45	{'min_samples_le
3	0.050727	0.001386	0.002415	0.000073	65	{'min_samples_le
4	0.050003	0.003110	0.002739	0.000548	85	{'min_samples_le

```

"""
# plotting accuracies with min_samples_leaf
plt.figure()
plt.plot(scores["param_min_samples_leaf"],
        scores["mean_train_score"],
        label="training accuracy")
plt.plot(scores["param_min_samples_leaf"],
        scores["mean_test_score"],
        label="test accuracy")
plt.xlabel("min_samples_leaf")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
"""

'\n# plotting accuracies with min_samples_leaf\nplt.figure()\nplt.plot(scores["param_min_samples_lea
f"], \n          scores["mean_train_score"], \n          label="training accuracy")\nplt.plot(scores["pa
ram_min_samples_leaf"], \n          scores["mean_test_score"], \n          label="test accuracy")\nplt.x
label("min samples leaf")\nplt.ylabel("Accuracy")\nplt.legend()\nplt.show()\n'

# GridSearchCV to find optimal min_samples_split
from sklearn.model_selection import KFold

```

```

from sklearn.model_selection import GridSearchCV


# specify number of folds for k-fold CV
n_folds = 5

# parameters to build the model on
parameters = {'min_samples_split': range(5, 200, 20)}

# instantiate the model
dtree = DecisionTreeClassifier(criterion = "gini",
                              random_state = 100)

# fit tree on training data
tree = GridSearchCV(dtree, parameters,
                    cv=n_folds,
                    scoring="accuracy")
tree.fit(X_train, y_train)

```



```

# scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()

```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_min_samples_split	params	split0_test_score	spli
0	0.075732	0.001588	0.003151	0.000410	5	{'min_samples_split': 5}	0.811982	
1	0.072139	0.002054	0.002695	0.000054	25	{'min_samples_split': 25}	0.825006	
2	0.069699	0.001631	0.002750	0.000079	45	{'min_samples_split': 45}	0.835188	
3	0.067109	0.001148	0.002617	0.000048	65	{'min_samples_split': 65}	0.839451	
4	0.066176	0.002326	0.002656	0.000063	85	{'min_samples_split': 85}	0.846081	

```

"""
# plotting accuracies with min_samples_leaf
plt.figure()
plt.plot(scores["param_min_samples_split"],
         scores["mean_train_score"],
         label="training accuracy")
plt.plot(scores["param_min_samples_split"],
         scores["mean_test_score"],
         label="test accuracy")
plt.xlabel("min_samples_split")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
"""

'\n# plotting accuracies with min_samples_leaf\nplt.figure()\nplt.plot(scores["param_min_samples_split"], \n         scores["mean_train_score"], \n         label="training accuracy")\nplt.plot(scores["param_min_samples_split"], \n         scores["mean_test_score"], \n         label="test accuracy")\nplt.xlabel("min_samples_split")\nplt.ylabel("Accuracy")\nplt.legend()\nplt.show()\n'

```

```

# Create the parameter grid
param_grid = {
    'max_depth': range(5, 15, 5),
    'min_samples_leaf': range(50, 150, 50),
    'min_samples_split': range(50, 150, 50),
    'criterion': ["entropy", "gini"]
}


n_folds = 5

# Instantiate the grid search model
dtree = DecisionTreeClassifier()
grid_search = GridSearchCV(estimator = dtree, param_grid = param_grid,
                           cv = n_folds, verbose = 1)

```

```
# Fit the grid search to the data
grid_search.fit(X_train,y_train)
```

Fitting 5 folds for each of 16 candidates, totalling 80 fits



```
# cv results
cv_results = pd.DataFrame(grid_search.cv_results_)
cv_results
```

3	0.032360	0.001393	0.002821	0.000553	entropy	5
4	0.051469	0.001100	0.002973	0.000313	entropy	10
5	0.050876	0.001384	0.002635	0.000130	entropy	10

```
# printing the optimal accuracy score and hyperparameters
print("best accuracy", grid_search.best_score_)
print(grid_search.best_estimator_)
```

```
best accuracy 0.8510400232064759
DecisionTreeClassifier(max_depth=10, min_samples_leaf=50, min_samples_split=50)
```

```
# model with optimal hyperparameters
clf_gini = DecisionTreeClassifier(criterion = "gini",
                                random_state = 100,
                                max_depth=10,
                                min_samples_leaf=50,
                                min_samples_split=50)
clf_gini.fit(X_train, y_train)
```


```
DecisionTreeClassifier
DecisionTreeClassifier(max_depth=10, min_samples_leaf=50, min_samples_split=50,
                      random_state=100)
```

```
# accuracy score
clf_gini.score(X_test,y_test)
```

```
0.850922753895458
```

```
# plotting the tree
dot_data = StringIO()
export_graphviz(clf_gini, out_file=dot_data,feature_names=features,filled=True,rounded=True)
```

```
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```



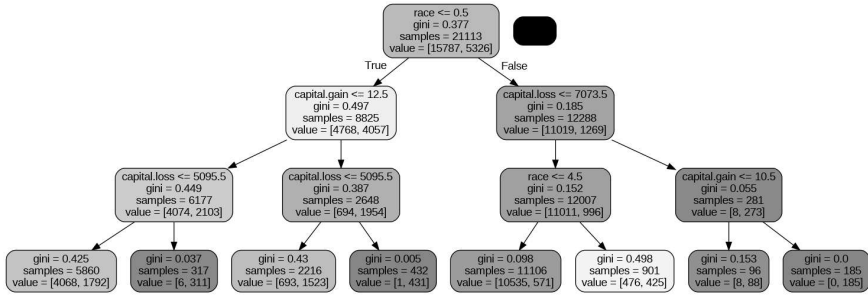
```
# tree with max_depth = 3
clf_gini = DecisionTreeClassifier(criterion = "gini",
                                random_state = 100,
                                max_depth=3,
                                min_samples_leaf=50,
                                min_samples_split=50)
clf_gini.fit(X_train, y_train)
```

```
# score
print(clf_gini.score(X_test,y_test))
```

```
0.8393192617968837
```

```
# plotting tree with max_depth=3
dot_data = StringIO()
export_graphviz(clf_gini, out_file=dot_data,feature_names=features,filled=True,rounded=True)
```

```
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```



```
# classification metrics
from sklearn.metrics import classification_report, confusion_matrix
y_pred = clf_gini.predict(X_test)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.96	0.90	6867
1	0.77	0.47	0.59	2182
accuracy			0.84	9049
macro avg	0.81	0.71	0.74	9049
weighted avg	0.83	0.84	0.82	9049

+ Code

+ Text

```
# confusion matrix
print(confusion_matrix(y_test,y_pred))

[[6564 303]
 [1151 1031]]
```



### **Conclusion:**

1. Discuss about the how categorical attributes have been dealt with during data pre-processing.

The Adult Census Income Dataset contains a mix of categorical and numerical attributes. Categorical attributes are those that can be divided into categories, such as "marital-status" and "race". Numerical attributes are those that can be represented as numbers, such as "age" and "hours-per-week".

Decision trees can handle both categorical and numerical attributes, but they require different approaches to deal with them effectively. Categorical attributes are typically encoded as dummy variables, which are binary variables that indicate the presence or absence of a particular category. For example, the "marital-status" attribute could be encoded as three dummy variables: "married", "single", and "divorced".

In this study, the categorical attributes were encoded as dummy variables using the `LabelEncoder()` class from the `sklearn.preprocessing` library. The numerical attributes were not pre-processed.

2. Discuss the hyper-parameter tuning done based on the decision tree obtained.

The following hyperparameters were tuned:

**max\_depth:** This is the maximum depth of the decision tree. A deeper tree will have more splits and can potentially learn more complex patterns, but it can also be more prone to overfitting.

**min\_samples\_split:** This is the minimum number of samples required to split a node. A higher value will prevent the tree from splitting too much and overfitting the training data.

**min\_samples\_leaf:** This is the minimum number of samples required in a leaf node. A higher value will prevent the tree from creating too many leaf nodes and underfitting the training data.

**criterion:** This is the splitting criterion used to determine the best split at each node. The most common criterion is Gini impurity, but other criteria such as entropy can also be used.





The hyperparameters were tuned using a grid search with cross-validation. This means that the model was trained and evaluated on different combinations of hyperparameters, and the best combination was selected.

The optimal hyperparameters were found to be:

```
max_depth = 3  
min_samples_split = 50  
min_samples_leaf = 50  
criterion = "gini"
```

3. Comment on the accuracy, confusion matrix, precision, recall and F1 score obtained.

**Accuracy:** The accuracy is the percentage of predictions that were correct. In this case, the accuracy is 83%. This means that the model correctly classified 83% of the test set samples.

**Precision:** Precision is the percentage of positive predictions that were actually positive. In this case, the precision is 85%. This means that 85% of the samples that the model predicted to be >50K were actually >50K.

**Recall:** Recall is the percentage of actual positives that were correctly predicted. In this case, the recall is 96%. This means that 96% of the samples that were actually >50K were correctly predicted by the model.

**F1 score:** The F1 score is a weighted average of precision and recall. In this case, the F1 score is 90%. This means that the model has a good balance of precision and recall